



CIS5560 Term Project Tutorial



Authors: Adnan Elahi; Jinhui Liu; Shilpa Konde Deshmukh; Siying Chen

Instructor: [Jongwook Woo](#)

Date: 05/18/2021

Lab Tutorial

Adnan Elahi (aelahi@calstatela.edu)

Jinhui Liu (jliu2@calstatela.edu)

Shilpa Konde Deshmukh (skonded@calstatela.edu)

Siying Chen (schen112@calstatela.edu)

05/18/2021

Fare Prediction for New York taxi

Objectives

The aim of this tutorial is to predict the fare by giving trip distance, trip time, and number of passengers. In this hands-on lab, you will learn how to:

- Get data manually
- AzureML Model
- Create Spark cluster
- Spark ML Model
- Visualization
- <https://gallery.cortanaintelligence.com/Experiment/Kaggle-Final>

- <https://gallery.cortanaintelligence.com/Experiment/FareTrip-Final>

Platform Spec

- Microsoft Azure Machine Learning
- Databricks Community Edition
- # of CPU cores: 1(AZML)/2(Databricks)
- # of nodes: 1
- Total Memory Size: 10GB(AZML)/15.3GB(Databricks)

Task 1: Get data manually and modify the data

This step is to get data manually, and prepare the datasets for azure ml and databricks

1. Download the first dataset from Kaggle: <https://www.kaggle.com/microize/newyork-yellow-taxi-trip-data-2020-2019> (we need the data from 2019/11 to 2020/02 here).
2. Download the second dataset from Chris Whong: https://chriswhong.com/open-data/foil_nyc_taxi/ (we need both Fare and Trip data).
3. Put all the Kaggle datasets in a new folder and combine the Kaggle datasets by following cmd code:

```
cat *.csv > KaggleTaxi.csv
```

4. Copy the first 113400 rows from KaggleTaxi.csv for a 10MB sample dataset by following cmd

code:

```
head -n 113400 KaggleTaxi.csv > KaggleTaxiSample.csv
```

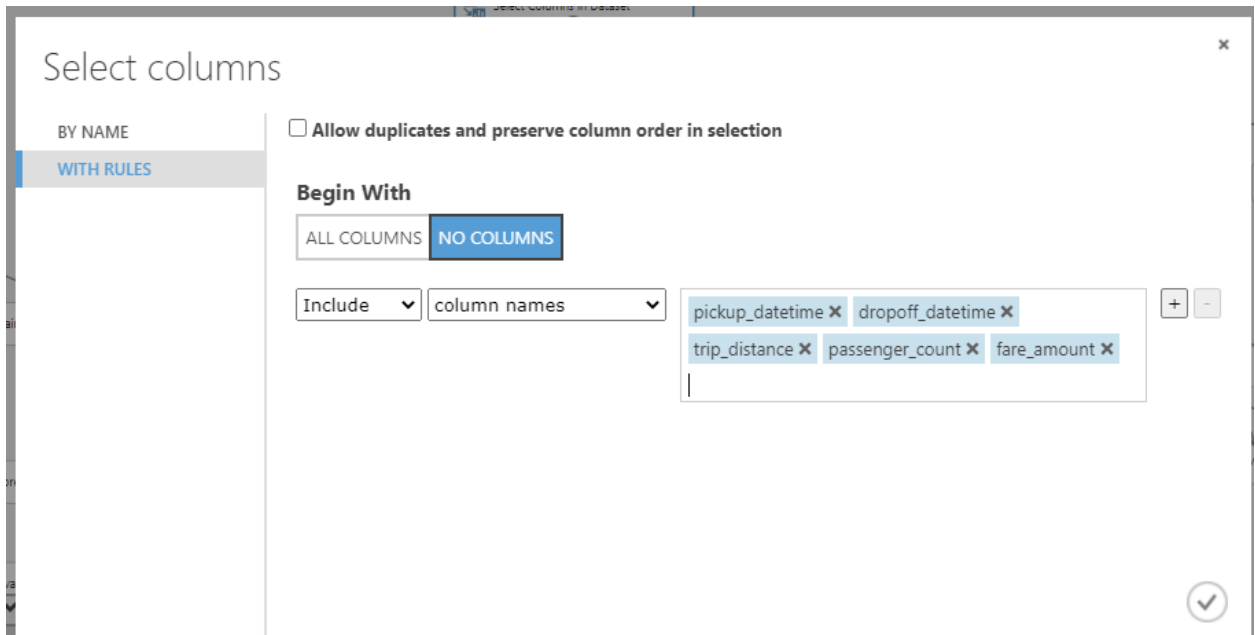
5. Do the same thing for fare and trip sample. 92530 rows for fareSample.csv and 62700 rows for tripSample.csv

Task 2: Azure Machine Learning

This step will introduce Azure ML

1. Open a browser and sign in using the Microsoft account associated with your Azure ML account at <https://studio.azureml.net/>
2. Click New at bottom left -> Dataset -> From Local File -> select **KaggleTaxiSample.csv** to upload the sample dataset to Azure.

3. Click New -> Experiment -> Blank Experiment and rename it to KaggleTaxi.
4. Click Saved Datasets on the left -> My Datasets and drag KaggleTaxiSample.csv to the canvas.
5. Search for “**Select Columns in Dataset**” and drag it to the canvas under Dataset module.
6. Link the Dataset output to the input of **Select Columns** in Dataset module.
7. Click **Select Columns** in Dataset module and click **Launch Column Selector**.
8. Choose **With Rules**, **No Columns** and select pickup_datetime, dropoff_datetime, trip_distance, passenger_count and fare_amount.



9. Search and drag **Split Data** to the canvas, connect it to the select columns module. And change the properties as in the picture.

Properties

Project

Split Data

Splitting mode

Split Rows

Fraction of rows in the first...

.7

☒ Randomized split

Random seed

12345

Stratified split

False

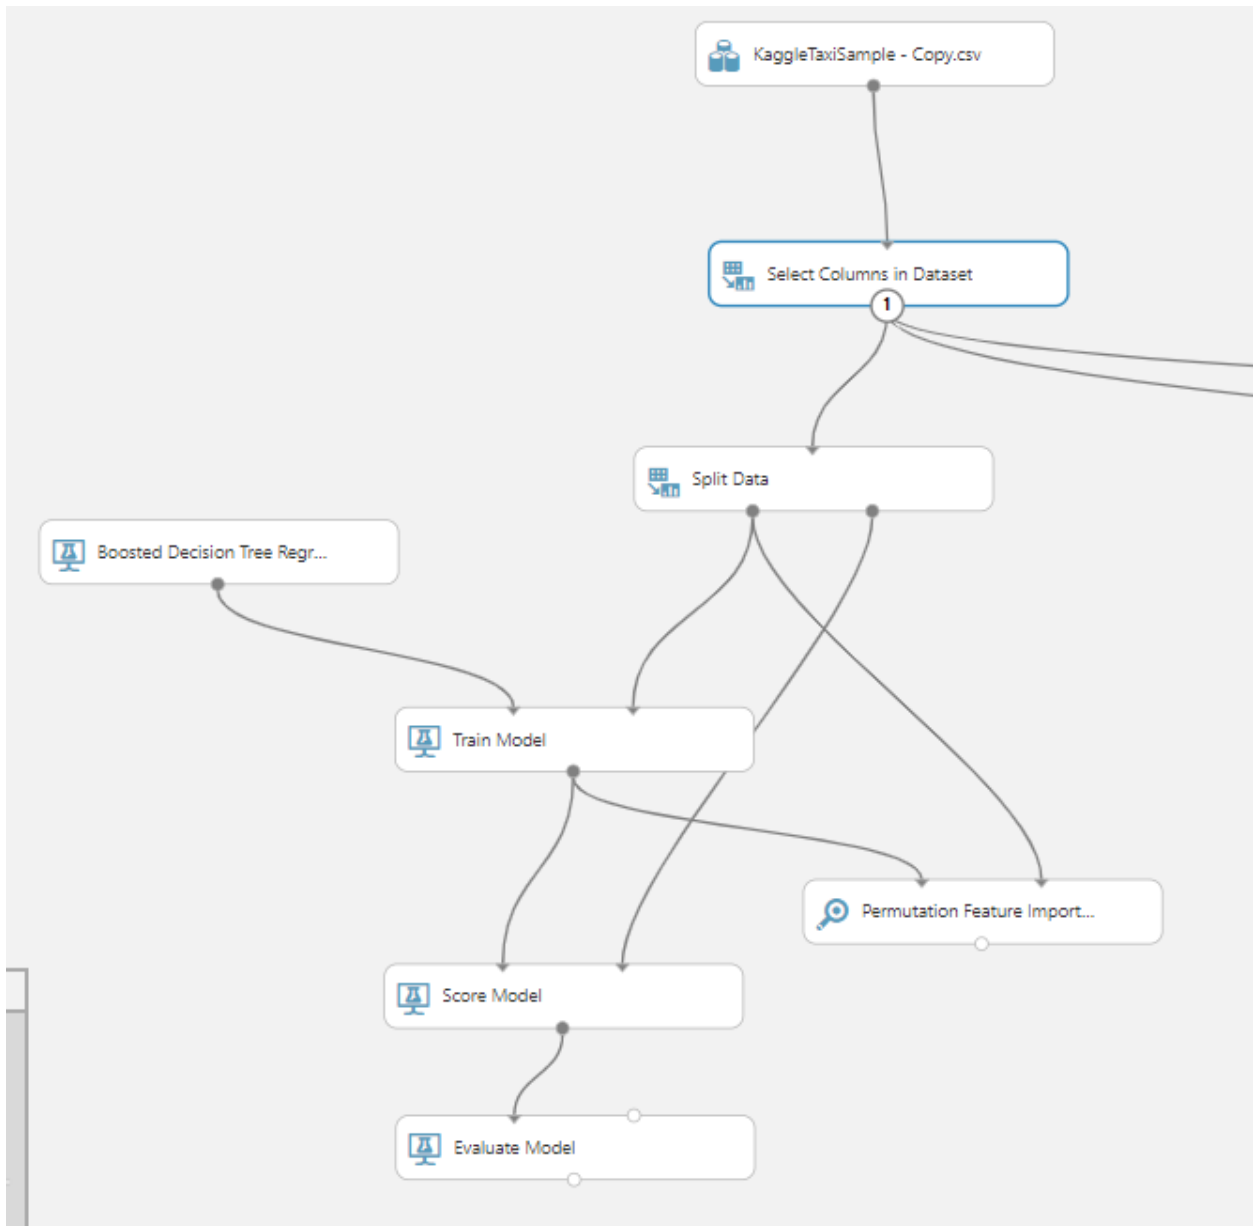
10. Search and drag **Train Model** under **Split Data**. Select “fare_amount” by using column selector. Link the Results dataset1(left) from Split Data to Dataset(right) of Train Model.

11. Search and drag **Boosted Decision Tree Regression**. Set up the properties as following: Create trainer mode -> Single Parameter. Maximum number of leaves per tree ->20. Minimum number of Samples per leaf node -> 10. Learning rate -> 0.2. Total number of trees constructed ->100. Random number seed ->12345. Allow unknown categorical levels -> checked. And link it to Untrained model(left) of Train Model.

12. Search and drag **Score Model** under **Train Model**. Link Trained Model(left) from Trained Model of Train Model, and Dataset(right) from Results Dataset2(right) from Split Data. Append score columns to output -> checked.

13. Search and drag **Permutation Feature Importance** next to **Score Model**. And link Trained Model(left) from Trained Model, Test Data(right) from Results dataset1 of Split Data. Set the properties as Random seed -> 12345, Metric for measuring performance -> Regression-Root Mean Squared Error.

14. Search and drag **Evaluate Model** under **Score Model**, link Scored data(left) from Scored datasets of Score Model. The general view should be like this:




15. Save and Run it.


16. Repeat from step 9 to 15 for **Linear Regression**. (Search and drag **Linear Regression** instead of **Boosted Decision Tree Regression**) and Setup the properties of Linear Regression as following: Solution method -> Ordinary Least Squares. L2 regularization weight -> 0.001. Include intercept term -> checked. Random number seed -> 12345. Allow unknown categorical levels -> checked.

17. Repeat from step 9 to 15 for **Decision Forest Regression**. And setup the properties as following:


Properties Project


Decision Forest Regression

Resampling method 


Bagging 

Create trainer mode


Single Parameter 

Number of decision trees 


8

Maximum depth of the decision trees 


32

Number of random splits per node 

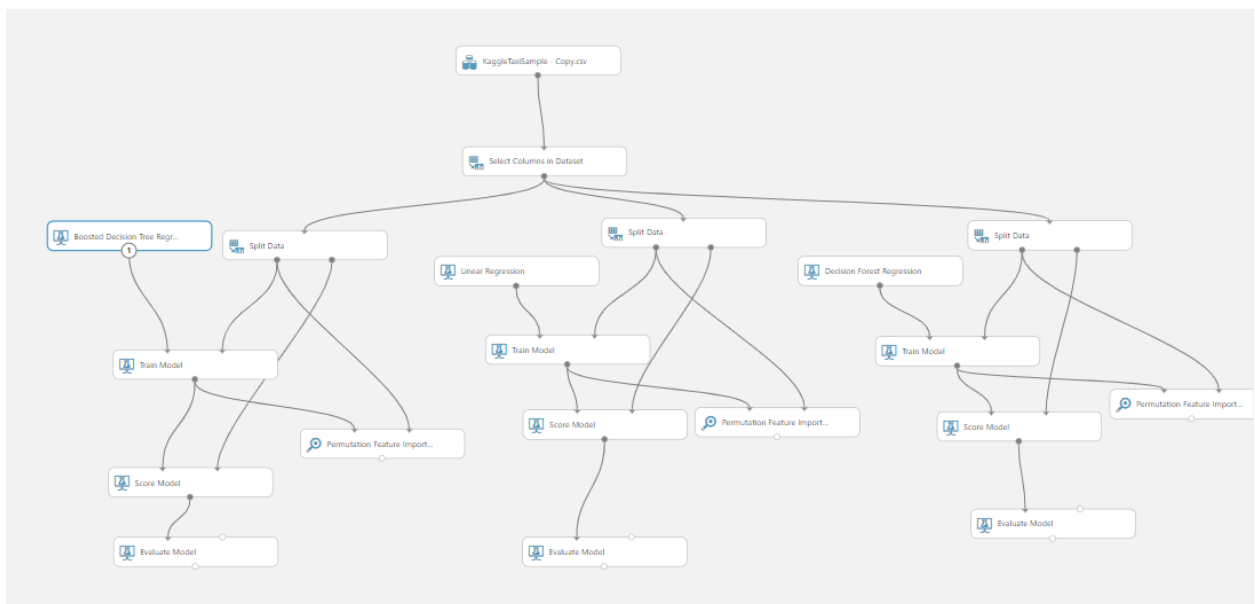
128

Minimum number of samples per leaf node 

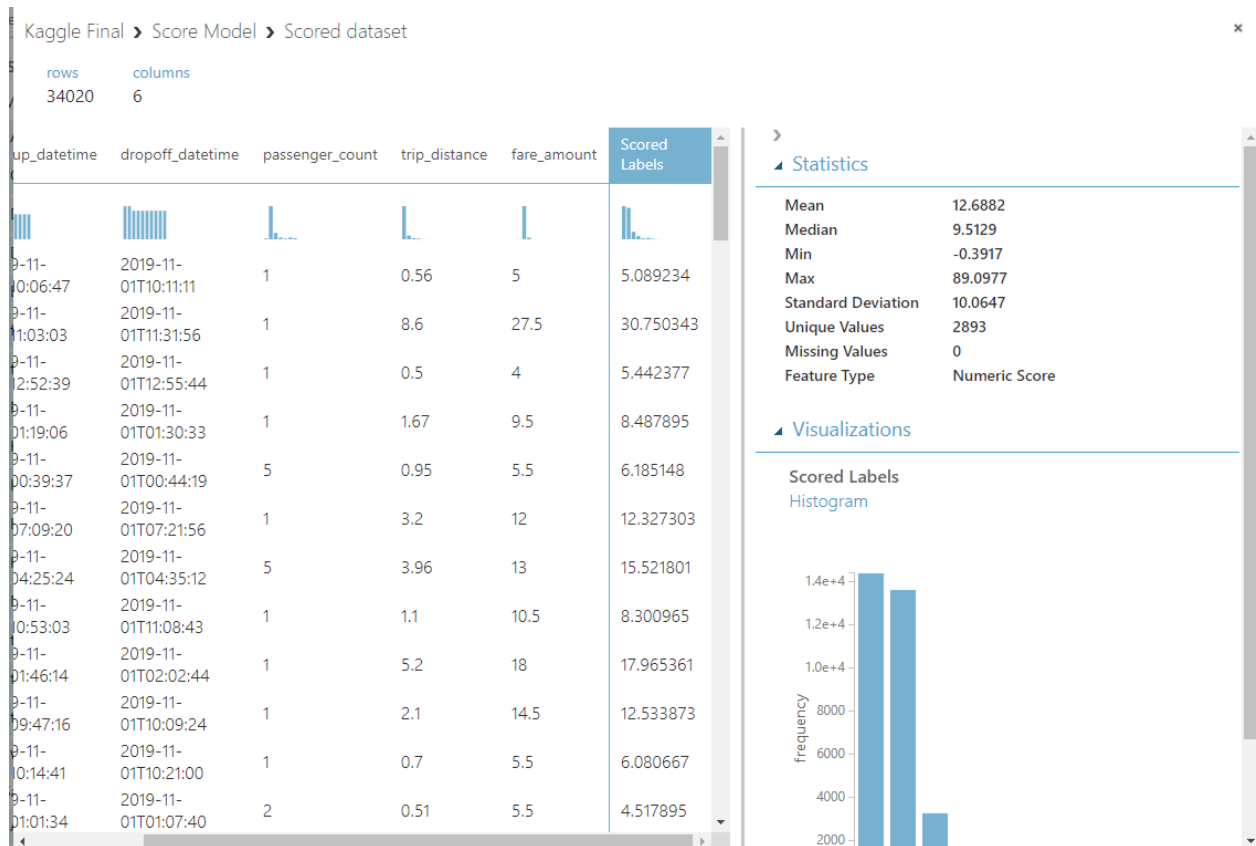
1

☒ Allow unknown values for categorical features 

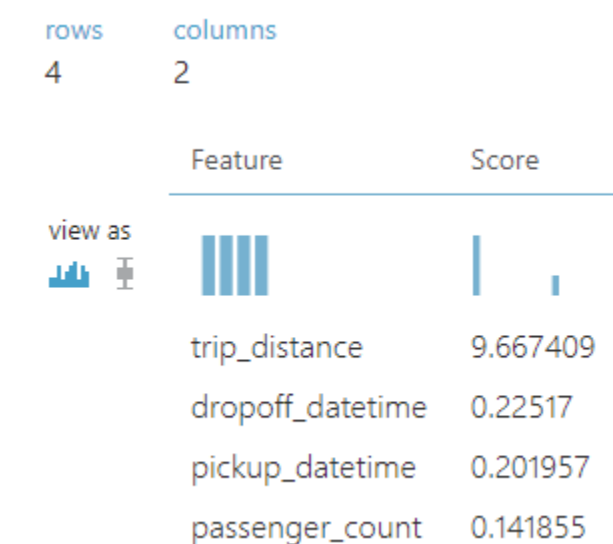
18. The final view will be like this:



19. Right click **Scored dataset** of **Score Model** to Visualize the output. The Score Labels are the predicted values.



20. Right click **Feature importance of Permutation Feature Importance** to visualize the importance of each features.




21. Right click **Evaluation Results** of **Evaluate Model** to visualize the **RMSE** and **Coefficient of Determination** or other evaluation results.

Metrics

Mean Absolute Error	1.990758
Root Mean Squared Error	5.552618
Relative Absolute Error	0.276118
Relative Squared Error	0.23958
Coefficient of Determination	0.76042

Error Histogram



22. Do the same thing from step 19 to 21 for all three algorithms and compare the RMSE and Coefficient of Determination of them. Think which algorithm is the best(Boosted Decision Tree Regression is the best in this case).

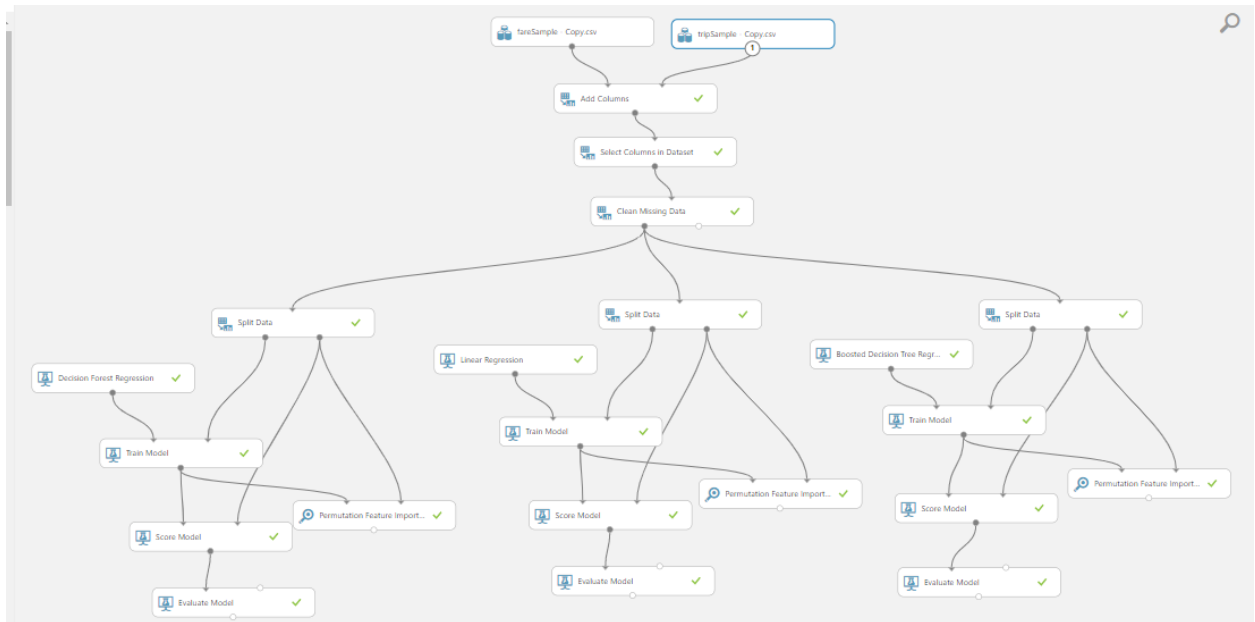
23. Upload **fareSample.csv** and **tripSample.csv** to Azure and drag them to a new canvas named as **FareTrip Final**.

24. Search and drag **Add Columns** and connect both datasets to it.

25. Drag **Select Columns** in Dataset under **Add Columns**. And select fare_amount, trip_time_in_secs, trip_distance and passenger_count by using column selector.

26. Drag **Clean Missing Data** under **Select Columns**, and set the properties as following: Selected columns -> All columns. Minimum missing value ratio -> 0. Maximum missing value ratio -> 1. Cleaning mode -> Remove entire row.








27. Repeat step 9 to 17 for all three algorithms, save and run the experiment.



28. Repeat step 19 to 22. to compare the results for all three algorithms and also compare the all the results from two different datasets/experiments.(following is results for BDTR of FareTrip Final)

rows
18810

columns
5


	fare_amount	passenger_count	trip_time_in_secs	trip_distance	Scored Labels
view as  					
7	1		480	1.08	7.193695
16	5		1140	3.96	16.346048
8.5	5		480	1.95	8.657975
10	2		660	2.37	10.10342
5	1		300	0.69	5.177501
9.5	5		720	1.58	9.379555
8	5		480	1.59	7.782753
5	5		120	1.08	4.828849
4.5	1		180	0.59	4.287382
7	1		480	1.24	7.228472
9	1		600	1.73	8.550353
4	1		120	0.66	4.179364
8	6		480	1.46	7.524806
13	1		1200	3.1	14.968396
4.5	1		180	0.71	4.46234
26	1		1260	8.35	26.235025
12.5	2		840	2.92	12.258604

Statistics

Mean	11.7125
Median	8.5648
Min	-0.0318
Max	83.6398
Standard Deviation	9.4398
Unique Values	3833
Missing Values	0
Feature Type	Numeric Score

Visualizations

Scored Labels Histogram



rows
3


columns
2

	Feature	Score
view as		
	trip_distance	1.314475
	trip_time_in_secs	0.397112
	passenger_count	0.003035

Metrics

Mean Absolute Error	0.671575
Root Mean Squared Error	3.4599
Relative Absolute Error	0.106469
Relative Squared Error	0.118126
Coefficient of Determination	0.881874

Error Histogram

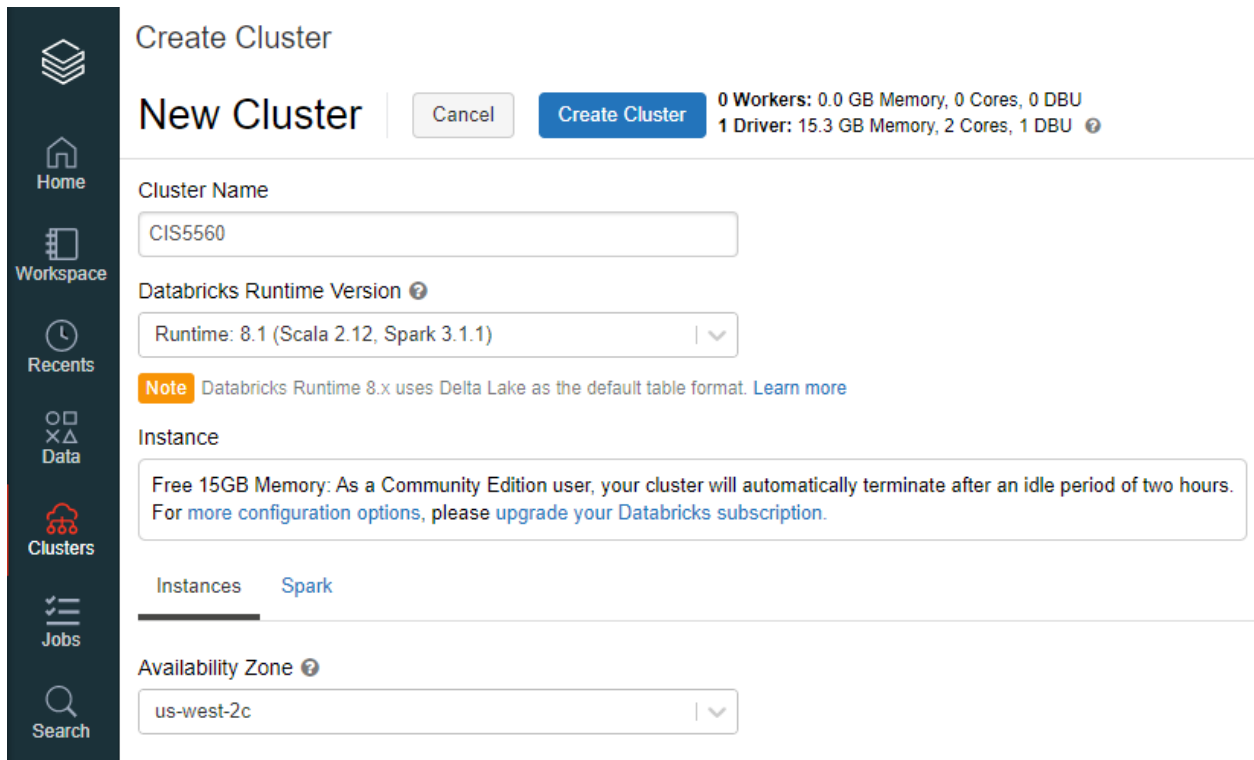


Task 3: Databricks

This step will introduce Databricks

1. Open a browser and sign in using the Databricks account at <https://community.cloud.databricks.com/>

2. Once sign in to Databricks community edition, click **Clusters** on the left side -> Create Cluster and name the Cluster Name as **CIS5560**. Choose Databricks Runtime Version as Runtime: **8.1 (Scala 2.12, Spark 3.1.1)**



Create Cluster

New Cluster

Cancel Create Cluster

0 Workers: 0.0 GB Memory, 0 Cores, 0 DBU
1 Driver: 15.3 GB Memory, 2 Cores, 1 DBU ⓘ

Cluster Name

CIS5560

Databricks Runtime Version ⓘ

Runtime: 8.1 (Scala 2.12, Spark 3.1.1) | v

Note Databricks Runtime 8.x uses Delta Lake as the default table format. [Learn more](#)


Instance

Free 15GB Memory: As a Community Edition user, your cluster will automatically terminate after an idle period of two hours. For [more configuration options](#), please [upgrade your Databricks subscription](#).

Instances Spark

Availability Zone ⓘ

us-west-2c | v

3. Click the symbol on the left top  to back the main page.
4. Click **Import & Explore Data** -> Drop files to upload, or click to browse to upload the KaggleTaxiSample.csv and choose **Create Table in Notebook**

Create New Table

Data source ⓘ

Upload File S3 DBFS Other Data Sources Partner Integrations

DBFS Target Directory ⓘ

/FileStore/tables/ (optional) Select

Files uploaded to DBFS are accessible by everyone who has access to this workspace. [Learn more](#)

Files ⓘ

KaggleTaxiSample-2.csv
10.5 MB
[Remove file](#)

✓ File uploaded to /FileStore/tables/KaggleTaxiSample-2.csv

Create Table with UI Create Table in Notebook ⓘ

5. Once finished uploading the dataset, click **Workspace** on the left side and open the dataset (if the dataset does not show up automatically) rename it as KTSample.

6. Click Workspace -> left click cis5560 folder if have -> right click on the blank space -> Create -> Notebook. Give the name as KaggleTSampleTest, default language as Python, Cluster using the CIS5560 which just created.

7. Insert following code for **apply packages**:

```
from pyspark.sql.types import *
from pyspark.sql.functions import *

from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession

from pyspark.ml import Pipeline
from pyspark.ml.regression import GBRegressor
from pyspark.ml.regression import LinearRegression
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
from pyspark.ml.evaluation import BinaryClassificationEvaluator,
RegressionEvaluator
```

```
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
from pyspark.ml.regression import DecisionTreeRegressor
```

8. Insert a new cell by enter following code:

```
IS_SPARK_SUBMIT_CLI = True

if IS_SPARK_SUBMIT_CLI:
    sc = SparkContext.getOrCreate()
    spark = SparkSession(sc)
```

9. Insert a new cell by enter following code for **load data**:

```
file_location = "/FileStore/tables/KTSample.csv"
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these
will be ignored.
df1 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)
```

10. Insert a new cell by enter following code for **load data** and check the first 5 lines of the data:

```
temp_table_name1 = "KTSample_csv"

df1.createOrReplaceTempView(temp_table_name1)
if IS_SPARK_SUBMIT_CLI:
    KTSample = spark.read.csv('/FileStore/tables/KTSample.csv',
inferSchema=True, header=True)
else:
    KTSample = sqlContext.sql("select * from KTSample_csv")

KTSample.show(5)
```

```

1 if IS_SPARK_SUBMIT_CLI:
2     KTSample = spark.read.csv('/FileStore/tables/KTSample.csv', inferSchema=True, header=True)
3 else:
4     KTSample = sqlContext.sql("select * from KTSample_csv")
5
6 KTSample.show(5)

```

▶ (3) Spark Jobs

▶ KTSample: pyspark.sql.dataframe.DataFrame = [VendorID: integer, tpep_pickup_datetime: string ... 16 more fields]

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	tip_amount	tolls_amount	improvement_surcharge	total_amount	congestion_surcharge
1	1	2019-11-01 00:30:41	2019-11-01 00:32:25	1	0.0	1	N	145	145		3.0	0.5	0.5	0.0	0.0	0.0	4.3	0.0
2	1	2019-11-01 00:34:01	2019-11-01 00:34:09	1	0.0	1	N	145	145		2.5	0.5	0.5	0.0	0.0	0.0	3.8	0.0
2	2	2019-11-01 00:41:59	2019-11-01 00:42:23	1	0.0	1	N	193	193		2.5	0.5	0.5	0.95	0.0	0.0	4.75	0.0
1	2	2019-11-01 00:02:39	2019-11-01 00:02:51	1	0.0	1	N	193	193		2.5	0.5	0.5	0.95	0.0	0.0	4.75	0.0
2	2	2019-11-01 00:18:30	2019-11-01 00:18:39	2	0.0	1	N	226	226		2.5	0.0	0.5	0.0	0.0	0.0	3.3	0.0

only showing top 5 rows

11. Insert a new cell by enter following code for **select data** and **calculate the trip time in sec**:

```

timediff=KTSample.select('tpep_pickup_datetime',
'tpep_dropoff_datetime','passenger_count','trip_distance',col('fare_
amount').alias('label'))

df2=timediff.withColumn('tpep_pickup_datetime',to_timestamp(col('tpe
p_pickup_datetime')))\
.withColumn('tpep_dropoff_datetime',
to_timestamp(col('tpep_dropoff_datetime')))\
.withColumn('trip_time_in_secs',col("tpep_dropoff_datetime").cast(
"long") - col('tpep_pickup_datetime').cast("long"))

data=df2.select('passenger_count','trip_distance','trip_time_in_secs
','label')

data.show(5)

```

```

1 data=df2.select('passenger_count','trip_distance','trip_time_in_secs','label')
2
3 data.show(5)

```

▶ (1) Spark Jobs

▶ data: pyspark.sql.dataframe.DataFrame = [passenger_count: integer, trip_distance: double ... 2 more fields]

```

+-----+-----+-----+-----+
|passenger_count|trip_distance|trip_time_in_secs|label|
+-----+-----+-----+-----+
|              1|          0.0|             104|  3.0|
|              1|          0.0|              8|  2.5|
|              1|          0.0|              24|  2.5|
|              1|          0.0|              12|  2.5|
|              2|          0.0|              9|  2.5|
+-----+-----+-----+-----+

```

only showing top 5 rows

Command took 0.50 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:52 on My Cluster

12. Insert a new cell to setup **train** and **test** datasets by following code:

```

splits = data.randomSplit([0.7, 0.3])
train = splits[0]
test = splits[1].withColumnRenamed("label", "trueLabel")

```

13. Insert a new cell to setup **assembler** for GBT-Regression:

```

assembler = VectorAssembler(inputCols = ['passenger_count',
'trip_time_in_secs', 'trip_distance'], outputCol="features")
gbt = GBTRegressor(labelCol="label")

```

14. Insert a new cell to setup **ParamGrid** and **CrossValidator**:

```

paramGrid = ParamGridBuilder()\
    .addGrid(gbt.maxDepth, [2, 3])\
    .addGrid(gbt.maxIter, [10, 20])\
    .build()

cv = CrossValidator(estimator=gbt, evaluator=RegressionEvaluator(),
estimatorParamMaps=paramGrid)

```

15. Insert a new cell to setup **pipeline** and **train** the model:

```

pipeline = Pipeline(stages=[assembler, cv])
pipelineModel = pipeline.fit(train)

```

16. Insert a new cell for **prediction**:

```

predictions = pipelineModel.transform(test)
predicted = predictions.select("features", "prediction", "trueLabel")

```



```
predicted.createOrReplaceTempView("regressionPredictions")
```

17. Insert a new cell for **RMSE**:

```
evaluator = RegressionEvaluator(labelCol="trueLabel",  
predictionCol="prediction", metricName="rmse")  
rmse = evaluator.evaluate(predictions)  
print( "Root Mean Square Error (RMSE) for GBT Regression :", rmse)
```

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")  
2 rmse = evaluator.evaluate(predictions)  
3 print( "Root Mean Square Error (RMSE) for GBT Regression :", rmse)
```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for GBT Regression : 5.5503173808683615

Command took 2.16 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:52 on My Cluster

18. Insert a new cell for **R2**(Coefficient of Determination):

```
evaluator = RegressionEvaluator(labelCol="trueLabel",  
predictionCol="prediction", metricName="r2")  
r2 = evaluator.evaluate(predictions)  
print( "Coefficient of Determination (R2) for GBT Regression :", r2)
```

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")  
2 r2 = evaluator.evaluate(predictions)  
3 print( "Coefficient of Determination (R2) for GBT Regression :", r2)
```

► (1) Spark Jobs

Coefficient of Determination (R2) for GBT Regression : 0.7651795550483502

Command took 1.47 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:52 on My Cluster

19. Now we need to setup **Linear Regression**. Insert the following code for **assembler, lr and pipeline1**:

```
assembler = VectorAssembler(inputCols = ['passenger_count',  
'trip_time_in_secs', 'trip_distance'], outputCol="features")  
lr = LinearRegression(labelCol="label", featuresCol="features",  
maxIter=10, regParam=0.3)  
pipeline1 = Pipeline(stages=[assembler, lr])
```

20. Code for **paramGrid1**:

```
paramGrid1 = ParamGridBuilder().addGrid(lr.regParam, [0.3,  
0.01]).addGrid(lr.maxIter, [10, 5]).build()  
trainval = TrainValidationSplit(estimator=pipeline1,  
evaluator=RegressionEvaluator(), estimatorParamMaps=paramGrid1,  
trainRatio=0.8)
```

21. Tran the model:

```
pipelineModel = trainval.fit(train)
```

22. Transform the model with test dataset:

```
predictions = pipelineModel.transform(test)
```

23. Setup for prediction:

```
predicted = predictions.select("features", "prediction", "trueLabel")
predicted.createOrReplaceTempView("regressionPredictions")
```

24. Setup RMSE for LR:

```
evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
print ("Root Mean Square Error (RMSE) for Linear Regression :", rmse)
```

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")
2 rmse = evaluator.evaluate(predictions)
3 print ("Root Mean Square Error (RMSE) for Linear Regression :", rmse)
```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for Linear Regression : 5.800560237034428

Command took 1.76 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:53 on My Cluster

25. Setup R2 for LR:

```
evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="r2")
r2 = evaluator.evaluate(predictions)
print( "Coefficient of Determination (R2) for Linear Regression :",
r2)
```

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")
2 r2 = evaluator.evaluate(predictions)
3 print( "Coefficient of Determination (R2) for Linear Regression :", r2)
```

► (1) Spark Jobs

Coefficient of Determination (R2) for Linear Regression : 0.7435278849894993

Command took 1.55 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:53 on My Cluster

26. Set up Decision Forest Regression by following code:

```
assembler = VectorAssembler(inputCols = ['passenger_count',
'trip_time_in_secs', 'trip_distance'], outputCol="features")
dt = DecisionTreeRegressor(labelCol="label", featuresCol="features")
paramGrid2 = ParamGridBuilder()\
```

```

        .addGrid(dt.maxDepth, [2,3])\
        .addGrid(dt.maxBins, [10,20])\
        .build()
dtcv = CrossValidator(estimator = dt, estimatorParamMaps = paramGrid2,
evaluator = RegressionEvaluator(), numFolds=2)
pipeline2 = Pipeline(stages=[assembler, dtcv])
pipelineModel = pipeline2.fit(train)
predictions = pipelineModel.transform(test)
predicted = predictions.select("features", "prediction", "trueLabel")
predicted.createOrReplaceTempView("regressionPredictions")

```

27. Setup RMSE:

```

evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
print( "Root Mean Square Error (RMSE) for Decision Forest
Regression :", rmse)

```

```

1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")
2 rmse = evaluator.evaluate(predictions)
3 print( "Root Mean Square Error (RMSE) for Decision Forest Regression :", rmse)

```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for Decision Forest Regression : 6.204744251541987

Command took 1.99 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:53 on My Cluster

28. Setup R2:

```

evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="r2")
r2 = evaluator.evaluate(predictions)
print( "Coefficient of Determination (R2) for Decision Forest
Regression :", r2)

```

```

1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")
2 r2 = evaluator.evaluate(predictions)
3 print( "Coefficient of Determination (R2) for Decision Forest Regression :", r2)

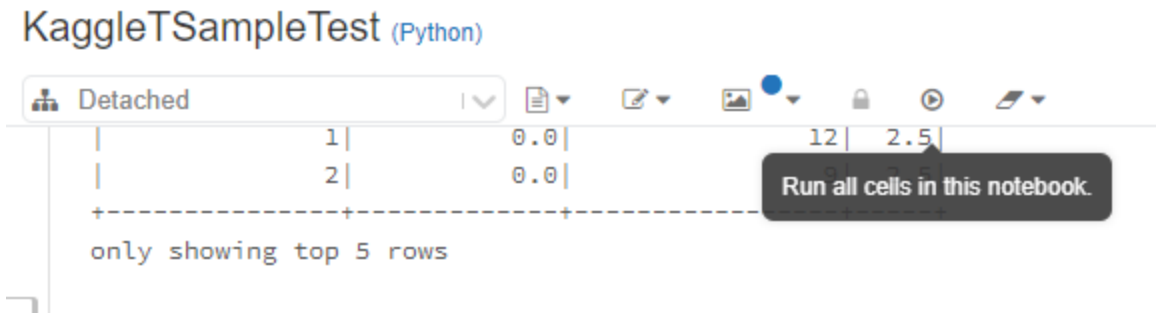
```

► (1) Spark Jobs

Coefficient of Determination (R2) for Decision Forest Regression : 0.7065405889317031

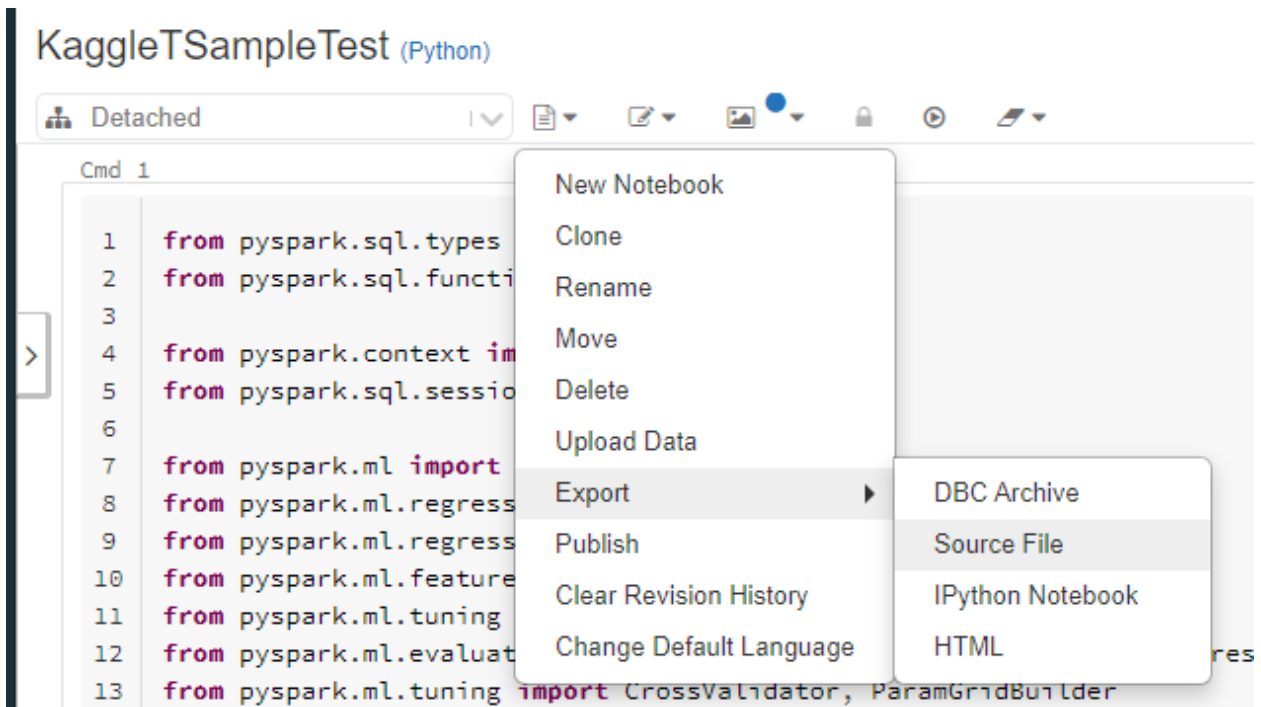
Command took 1.78 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:44:53 on My Cluster

29. Click Run all to run all the codes



30. Now you can compare the results for all three different models. For the **paramGrid** in GBT-Regression, we using **[2,3]** for **maxDepth** and **[10,20]** for **maxIter**. You may want to using **[2,5]** for **maxDepth** and **[10,100]** for **maxIter** to have more accuracy results.

31. You may also export the code as a **py** file for **Hadoop Spark** and **ipynb** file for **IPython Notebook** as shown in the picture:



32. Now repeat step 3 to 6 for upload fareSample.csv and tripSample.csv, renamed them as **fareSample** and **tripSample**. And create a new notebook named **FareTripSampleTest**.

33. Repeat step 7 and 8 to setup the environment:

```

1  from pyspark.sql.types import *
2  from pyspark.sql.functions import *
3
4  from pyspark.context import SparkContext
5  from pyspark.sql.session import SparkSession
6
7  from pyspark.ml import Pipeline
8  from pyspark.ml.regression import GBRegressor
9  from pyspark.ml.regression import LinearRegression
10 from pyspark.ml.feature import VectorAssembler
11 from pyspark.ml.tuning import ParamGridBuilder, TrainValidationSplit
12 from pyspark.ml.evaluation import BinaryClassificationEvaluator, RegressionEvaluator
13 from pyspark.ml.tuning import CrossValidator, ParamGridBuilder
14 from pyspark.ml.regression import DecisionTreeRegressor
15

```



Command took 0.92 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

Cmd 2

```

1  IS_SPARK_SUBMIT_CLI = True
2
3  if IS_SPARK_SUBMIT_CLI:
4      sc = SparkContext.getOrCreate()
5      spark = SparkSession(sc)

```

34. Load data by following codes:

```

file_location = "/FileStore/tables/fareSample.csv"
file_type = "csv"

```

```

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

```

The applied options are for CSV files. For other file types, these will be ignored.

```

df1 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

```

```

file_location = "/FileStore/tables/tripSample.csv"
file_type = "csv"

```

```

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these
will be ignored.
df2 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

temp_table_name1 = "fareSample_csv"
df1.createOrReplaceTempView(temp_table_name1)

temp_table_name2 = "tripSample_csv"
df2.createOrReplaceTempView(temp_table_name2)

if IS_SPARK_SUBMIT_CLI:
    fareSample = spark.read.csv('/FileStore/tables/fareSample.csv',
    inferSchema=True, header=True)
else:
    fareSample = sqlContext.sql("select * from fareSample_csv")
if IS_SPARK_SUBMIT_CLI:
    tripSample = spark.read.csv('/FileStore/tables/tripSample.csv',
    inferSchema=True, header=True)
else:
    tripSample = sqlContext.sql("select * from tripSample_csv")

```

35. Now we need to **combine** the two datasets by insert an **id** column to each dataset and **inner join** them:

```

fareSample = fareSample.withColumn('id',
monotonically_increasing_id())
tripSample = tripSample.withColumn('id',
monotonically_increasing_id())
data=fareSample.join(tripSample,"id")

```

36. Select the data, **clean** the **null** rows and show first **5** lines:

```

data1 = data.select('passenger_count', 'trip_time_in_secs',
'trip_distance', col('fare_amount').alias('label'))
dataaf = data1.dropna(how = 'any')
dataaf.show(5)

```

```
1 dataaf.show(5)
```

► (1) Spark Jobs

```
+-----+-----+-----+-----+
|passenger_count|trip_time_in_secs|trip_distance|label|
+-----+-----+-----+-----+
|          4|          382|          1.0|  6.5|
|          1|          259|          1.5|  6.0|
|          1|          282|          1.1|  5.5|
|          2|          244|          0.7|  5.0|
|          1|          560|          2.1|  9.5|
+-----+-----+-----+-----+
only showing top 5 rows
```

37. Setup **train** and **test** datasets:

```
splits = dataaf.randomSplit([0.7, 0.3])
train = splits[0]
test = splits[1].withColumnRenamed("label", "trueLabel")
```

38. Now we need to setup **GBT-Regression** which will be using in KaggleTaxi Sample by following codes:

```
assembler = VectorAssembler(inputCols = ['passenger_count',
'trip_time_in_secs', 'trip_distance'], outputCol="features")
gbt = GBTRegressor(labelCol="label")
paramGrid = ParamGridBuilder()\
    .addGrid(gbt.maxDepth, [2, 3])\
    .addGrid(gbt.maxIter, [10, 20])\
    .build()

cv = CrossValidator(estimator=gbt, evaluator=RegressionEvaluator(),
estimatorParamMaps=paramGrid)

pipeline = Pipeline(stages=[assembler, cv])
pipelineModel = pipeline.fit(train)
predictions = pipelineModel.transform(test)
predicted = predictions.select("features", "prediction", "trueLabel")
predicted.createOrReplaceTempView("regressionPredictions")
```

39. Setup **RMSE** and **R2**:

```
evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
print( "Root Mean Square Error (RMSE) for GBT Regression :", rmse)
```

```
evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="r2")
r2 = evaluator.evaluate(predictions)
print( "Coefficient of Determination (R2) for GBT Regression :", r2)
```

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")
2 rmse = evaluator.evaluate(predictions)
3 print( "Root Mean Square Error (RMSE) for GBT Regression :", rmse)
```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for GBT Regression : 10.205602402118904

Command took 2.58 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

Cmd 30

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")
2 r2 = evaluator.evaluate(predictions)
3 print( "Coefficient of Determination (R2) for GBT Regression :", r2)
```

► (1) Spark Jobs

Coefficient of Determination (R2) for GBT Regression : 0.13425932181419653

Command took 2.06 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

40. Setup Linear Regression, RMSE and R2 for LR:

```
assembler = VectorAssembler(inputCols = ['passenger_count',
'trip_time_in_secs', 'trip_distance'], outputCol="features")
```

```
lr = LinearRegression(labelCol="label", featuresCol="features",
maxIter=10, regParam=0.3)
```

```
pipeline1 = Pipeline(stages=[assembler, lr])
paramGrid1 = ParamGridBuilder().addGrid(lr.regParam, [0.3,
0.01]).addGrid(lr.maxIter, [10, 5]).build()
```

```
trainval = TrainValidationSplit(estimator=pipeline1,
evaluator=RegressionEvaluator(), estimatorParamMaps=paramGrid1,
trainRatio=0.8)
```

```
pipelineModel = trainval.fit(train)
predictions = pipelineModel.transform(test)
predicted = predictions.select("features", "prediction", "trueLabel")
predicted.createOrReplaceTempView("regressionPredictions")
```

```
evaluator = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
print( "Root Mean Square Error (RMSE) for Linear Regression :", rmse)
```

```
evaluator1 = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="r2")
r2 = evaluator1.evaluate(predictions)
```



```
print( "Coefficient of Determination (R2) for Linear Regression :",  
r2)
```

RMSE/R2 for Linear Regression

Cmd 40

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")  
2 rmse = evaluator.evaluate(predictions)  
3 print( "Root Mean Square Error (RMSE) for Linear Regression :", rmse)
```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for Linear Regression : 10.279668474240642

Command took 2.37 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

Cmd 41

```
1 evaluator1 = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")  
2 r2 = evaluator1.evaluate(predictions)  
3 print( "Coefficient of Determination (R2) for Linear Regression :", r2)
```

► (1) Spark Jobs

Coefficient of Determination (R2) for Linear Regression : 0.12164768198991038

Command took 1.86 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

41. Setup Decision Forest Regression, RMSE and R2 as before by following codes:

```
assembler = VectorAssembler(inputCols = ['passenger_count',  
'trip_time_in_secs', 'trip_distance'], outputCol="features")  
dt = DecisionTreeRegressor(labelCol="label", featuresCol="features")
```

```
paramGrid2 = ParamGridBuilder()\  
    .addGrid(dt.maxDepth, [2,3])\  
    .addGrid(dt.maxBins, [10,20])\  
    .build()
```

```
dtcv = CrossValidator(estimator = dt, estimatorParamMaps = paramGrid2,  
evaluator = RegressionEvaluator(), numFolds=2)
```

```
pipeline2 = Pipeline(stages=[assembler, dtcv])  
pipelineModel = pipeline2.fit(train)  
predictions = pipelineModel.transform(test)  
predicted = predictions.select("features", "prediction", "truelabel")  
predicted.createOrReplaceTempView("regressionPredictions")
```

```
evaluator = RegressionEvaluator(labelCol="trueLabel",  
predictionCol="prediction", metricName="rmse")  
rmse = evaluator.evaluate(predictions)  
print( "Root Mean Square Error (RMSE) for Decision Forest  
Regression :", rmse)
```

```
evaluator1 = RegressionEvaluator(labelCol="trueLabel",
predictionCol="prediction", metricName="r2")
r2 = evaluator1.evaluate(predictions)
print( "Coefficient of Determination (R2) for Decision Forest
Regression :", r2)
```

RMSE/R2 for Decision Forest Regression

Cmd 52

```
1 evaluator = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="rmse")
2 rmse = evaluator.evaluate(predictions)
3 print( "Root Mean Square Error (RMSE) for Decision Forest Regression :", rmse)
```

► (1) Spark Jobs

Root Mean Square Error (RMSE) for Decision Forest Regression : 10.290313716291786

Command took 2.19 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

Cmd 53

```
1 evaluator1 = RegressionEvaluator(labelCol="trueLabel", predictionCol="prediction", metricName="r2")
2 r2 = evaluator1.evaluate(predictions)
3 print( "Coefficient of Determination (R2) for Decision Forest Regression :", r2)
```

► (1) Spark Jobs

Coefficient of Determination (R2) for Decision Forest Regression : 0.11982756211888457

Command took 1.89 seconds -- by jliu2@calstatela.edu at 2021/5/14下午7:28:55 on My Cluster

42. Run all the code and now you can compare the results for all three different models. For the **paramGrid** in GBT-Regression, we using **[2,3]** for **maxDepth** and **[10,20]** for **maxIter**. You may want to using **[2,5]** for **maxDepth** and **[10,100]** for **maxIter** to have more accuracy results.

43. You may also export the code as a **py** file for **Hadoop Spark** and **ipynb** file for **IPythonNotebook**.

44. As we can see here, the GBT-Regression is the best model for both datasets based on the RMSE and R2. However, the Fare and Trip Samples return a different R2 compare to KaggleTaxi Sample. The reason could be the way how we combine the two(fare and trip) datasets into one for train and test.

Task 4: Hadoop Spark(optional)

This step will introduce Hadoop Spark

1. Open a Git Bash window. And upload the sample datasets, py file to Hadoop(you need to do this for three sample datasets and two py files. And using your user name instead of jliu2):

```
scp FareTripSampleTest.py jliu2@220.116.230.22:~/
```

```

LiuJH215@LiuJH215 MINGW64 ~/Desktop
$ scp FareTripSampleTest.py jliu2@220.116.230.22:~/
jliu2@220.116.230.22's password:
FareTripSampleTest.py                                100% 8215    44.7KB/s   00:00

LiuJH215@LiuJH215 MINGW64 ~/Desktop
$ scp KaggleTSampleTest.py jliu2@220.116.230.22:~/
jliu2@220.116.230.22's password:
KaggleTSampleTest.py                                100% 7472    42.7KB/s   00:00

```

2. Once finish upload, open a new Git Bush window and login into Hadoop by following code using your username and password:

```
ssh jliu2@220.116.230.22
```

3. using `ls -al` to check all the files uploaded correctly.

```

-rw-r--r-- 1 jliu2 jliu2 10405513 5월 7 15:20 tripSample.csv
-bash-4.2$ ls -al
한 계 4432640
drwx---w-. 6 jliu2 jliu2      4096 5월 18 11:46 .
drwxr-xr-x. 34 root  root      4096 3월 22 06:34 ..
-rw-----w-. 1 jliu2 jliu2    37534 5월 10 11:22 .bash_history
drwxr-xr-x. 2 root  root        40 3월 22 06:34 .beeline
drwxrwxrwx. 3 jliu2 jliu2      18 4월 6 02:34 .cache
drwxrwxrwx. 3 jliu2 jliu2      18 4월 6 02:34 .config
drwxrw--w-. 3 jliu2 jliu2      19 4월 6 03:07 .pkg
-rw-----w-. 1 jliu2 jliu2    7674 5월 18 11:45 .viminfo
-rw-r--r--. 1 jliu2 jliu2      8215 5월 18 11:45 FareTripSampleTest.py
-rw-r--r--. 1 jliu2 jliu2 10506769 5월 7 19:10 KTSample.csv
-rw-r--r--. 1 jliu2 jliu2      7472 5월 18 11:46 KaggleTSampleTest.py
-rw-r--r--. 1 jliu2 jliu2 2435115956 5월 8 07:35 KaggleTaxi.csv
-rw-r--r--. 1 jliu2 jliu2 842007423 5월 9 21:53 fare.csv
-rw-r--r--. 1 jliu2 jliu2 10480920 5월 7 15:16 fareSample.csv
-rw-r--r--. 1 jliu2 jliu2      6615 5월 10 06:16 prosample.py
-rw-r--r--. 1 jliu2 jliu2 1230406357 5월 9 22:08 trip.csv
-rw-r--r--. 1 jliu2 jliu2 10405513 5월 7 15:20 tripSample.csv
-bash-4.2$

```

4. Using following code to put all the dataset to **HDFS** and give permission to read/write:

```
hdfs dfs -put KTSample.csv
```

```
hdfs dfs -chmod -R o+w .
```

5. edit KaggleTSampleTest.py by using `vi`:

```
vi KaggleTSampleTest.py
```

6. Change file location and save it(change **jliu2** to your own user name).

```

file_location = "/user/jliu2/KTSample.csv"
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these will be ignored.
df1 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

# COMMAND -----

temp_table_name1 = "KTSample_csv"

df1.createOrReplaceTempView(temp_table_name1)

# COMMAND -----

if IS_SPARK_SUBMIT_CLI:
    KTSample = spark.read.csv('/user/jliu2/KTSample.csv', inferSchema=True, header=True)
else:
    KTSample = sqlContext.sql("select * from KTSample_csv")

```

7. run the py file by using following code:

```
spark-submit KaggleTSampleTest.py
```

8. After finished, you will see the **RMSE** and **R2** like the following, and you may also want to find the **RMSE** and **R2** for all **three** models, or you can modify the py file to display all results at the end.

```

21/05/18 12:32:36 INFO BlockManagerInfo: Added broadcast_2572_piece0 in memory on bigdata3.iscu.ac.kr:35644 (size: 32.9 KB, free:
365.8 MB)
21/05/18 12:32:36 INFO BlockManagerInfo: Added broadcast_2572_piece0 in memory on bigdata3.iscu.ac.kr:38712 (size: 32.9 KB, free:
365.8 MB)
21/05/18 12:32:36 INFO TaskSetManager: Finished task 1.0 in stage 1988.0 (TID 3733) in 494 ms on bigdata3.iscu.ac.kr (executor 1)
(1/2)
21/05/18 12:32:37 INFO TaskSetManager: Finished task 0.0 in stage 1988.0 (TID 3732) in 1204 ms on bigdata3.iscu.ac.kr (executor 2)
(2/2)
21/05/18 12:32:37 INFO YarnScheduler: Removed TaskSet 1988.0, whose tasks have all completed, from pool
21/05/18 12:32:37 INFO DAGScheduler: ResultStage 1988 (treeAggregate at RegressionMetrics.scala:57) finished in 1.214 s
21/05/18 12:32:37 INFO DAGScheduler: Job 1241 finished: treeAggregate at RegressionMetrics.scala:57, took 1.216544 s
('Coefficient of Determination (R2) for Decision Forest Regression :', 0.6963051700005948)
21/05/18 12:32:37 INFO SparkContext: Invoking stop() from shutdown hook
21/05/18 12:32:37 INFO AbstractConnector: Stopped Spark@5946d9bc{HTTP/1.1,[http/1.1]}{0.0.0.0:4040}
21/05/18 12:32:37 INFO SparkUI: Stopped Spark web UI at http://bigdata2.iscu.ac.kr:4040
21/05/18 12:32:45 INFO YarnClientSchedulerBackend: Interrupting monitor thread
21/05/18 12:32:45 INFO YarnClientSchedulerBackend: Shutting down all executors
21/05/18 12:32:45 INFO YarnSchedulerBackend$YarnDriverEndpoint: Asking each executor to shut down
21/05/18 12:32:45 INFO SchedulerExtensionServices: Stopping SchedulerExtensionServices
(serviceOption=None,
services=List(),
started=false)
21/05/18 12:32:45 INFO YarnClientSchedulerBackend: Stopped
21/05/18 12:32:45 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
21/05/18 12:32:45 INFO MemoryStore: MemoryStore cleared
21/05/18 12:32:45 INFO BlockManager: BlockManager stopped
21/05/18 12:32:45 INFO BlockManagerMaster: BlockManagerMaster stopped
21/05/18 12:32:45 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
21/05/18 12:32:45 INFO SparkContext: Successfully stopped SparkContext
21/05/18 12:32:45 INFO ShutdownHookManager: Shutdown hook called
21/05/18 12:32:45 INFO ShutdownHookManager: Deleting directory /tmp/spark-9431110e-7a6c-4f8f-8cab-ad6b9d99e35d/pyspark-bc5521d8-d
62c-47ee-9f0a-cb9c8b9d1d86
21/05/18 12:32:45 INFO ShutdownHookManager: Deleting directory /tmp/spark-491f0450-294e-439e-9b13-803febe74b3e
21/05/18 12:32:45 INFO ShutdownHookManager: Deleting directory /tmp/spark-9431110e-7a6c-4f8f-8cab-ad6b9d99e35d
-bash-4.2$ ^C
-bash-4.2$

```

9. Repeat step 5 to 7 for **FareTripSampleTest.py**

```
# COMMAND -----

file_location = "/user/jliu2/fareSample.csv"
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these will be ignored.
df1 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

file_location = "/user/jliu2/tripSample.csv"
file_type = "csv"

# CSV options
infer_schema = "true"
first_row_is_header = "true"
delimiter = ","

# The applied options are for CSV files. For other file types, these will be ignored.
df2 = spark.read.format(file_type) \
    .option("inferSchema", infer_schema) \
    .option("header", first_row_is_header) \
    .option("sep", delimiter) \
    .load(file_location)

# COMMAND -----

temp_table_name1 = "fareSample_csv"

df1.createOrReplaceTempView(temp_table_name1)

# COMMAND -----

temp_table_name2 = "tripSample_csv"

df2.createOrReplaceTempView(temp_table_name2)

# COMMAND -----

if IS_SPARK_SUBMIT_CLI:
    fareSample = spark.read.csv('/user/jliu2/fareSample.csv', inferSchema=True, header=True)
else:
    fareSample = sqlContext.sql("select * from fareSample_csv")

# COMMAND -----

if IS_SPARK_SUBMIT_CLI:
    tripSample = spark.read.csv('/user/jliu2/tripSample.csv', inferSchema=True, header=True)
else:
    tripSample = sqlContext.sql("select * from tripSample_csv")
```

10. After finished, you will see the **RMSE** and **R2** like the following, and you may also want to find the **RMSE** and **R2** for all **three** models, or you can modify the py file to display all results at the end.

```
21/05/18 13:32:28 INFO SparkContext: Created broadcast 2681 from broadcast at DAGScheduler.scala:1039
21/05/18 13:32:28 INFO DAGScheduler: Submitting 2 missing tasks from ResultStage 2025 (MapPartitionsRDD[4798] at treeAggregate at RegressionMetrics.scala:1039)
21/05/18 13:32:28 INFO YarnScheduler: Adding task set 2025.0 with 2 tasks
21/05/18 13:32:28 INFO TaskSetManager: Starting task 0.0 in stage 2025.0 (TID 3805, bigdata3.iscu.ac.kr, executor 1, partition 0, NODE_LOCAL, 832)
21/05/18 13:32:28 INFO TaskSetManager: Starting task 1.0 in stage 2025.0 (TID 3806, bigdata3.iscu.ac.kr, executor 2, partition 1, NODE_LOCAL, 832)
21/05/18 13:32:28 INFO BlockManagerInfo: Added broadcast_2681_piece0 in memory on bigdata3.iscu.ac.kr:39727 (size: 21.5 KB, free: 351.1 MB)
21/05/18 13:32:28 INFO BlockManagerInfo: Added broadcast_2679_piece0 in memory on bigdata3.iscu.ac.kr:44018 (size: 1640.8 KB, free: 351.2 MB)
21/05/18 13:32:28 INFO BlockManagerInfo: Added broadcast_2680_piece0 in memory on bigdata3.iscu.ac.kr:39727 (size: 33.0 KB, free: 349.5 MB)
21/05/18 13:32:28 INFO BlockManagerInfo: Added broadcast_2680_piece0 in memory on bigdata3.iscu.ac.kr:44018 (size: 33.0 KB, free: 351.2 MB)
21/05/18 13:32:29 INFO TaskSetManager: Finished task 1.0 in stage 2025.0 (TID 3806) in 430 ms on bigdata3.iscu.ac.kr (executor 2) (1/2)
21/05/18 13:32:29 INFO TaskSetManager: Finished task 0.0 in stage 2025.0 (TID 3805) in 821 ms on bigdata3.iscu.ac.kr (executor 1) (2/2)
21/05/18 13:32:29 INFO YarnScheduler: Removed TaskSet 2025.0, whose tasks have all completed, from pool
21/05/18 13:32:29 INFO DAGScheduler: ResultStage 2025 (treeAggregate at RegressionMetrics.scala:57) finished in 0.826 s
21/05/18 13:32:29 INFO DAGScheduler: Job 1278 finished: treeAggregate at RegressionMetrics.scala:57, took 0.828212 s
(Coefficient of Determination (R2) for Decision Forest Regression :, 0.3599860653814684)
```

11. According to all the results, the GBT Regression will give the best prediction. And compare to the databricks, R2 in Hadoop increased.

12. (Optional) you can also try the original dataset for Kaggle Taxi but be aware, it will take more than 6 hours to return the results.

References

1. **URL of Data Source:** <https://www.kaggle.com/microize/newyork-yellow-taxi-trip-data-2020-2019> and https://chriswhong.com/open-data/foil_nyc_taxi/
2. **GitHub URL:** <https://github.com/liujh215/CIS5560.git>
3. **Other Reference:**
 - <https://sparkbyexamples.com/spark/spark-difference-between-two-timestamps-in-seconds-minutes-and-hours/>
 - <https://spark.apache.org/docs/latest/mllib-decision-tree.html>
 - <https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.sql.DataFrame.dropna.html>
 - <https://spark.apache.org/docs/2.1.1/api/R/spark.gbt.html>
 - <https://gist.github.com/colbyford/daa4508f6d8d94a405e7bd3a50c5ed77>
 - https://docs.rapidminer.com/latest/studio/operators/modeling/predictive/trees/gradient_boosted_trees.html
 - https://docs.azuredatabricks.net/_static/notebooks/gbt-regression.html